

Progress Report: Identity Recognition With Masks

Group member: Chen Zhao, Ruihua Qiao, Zijian Chen

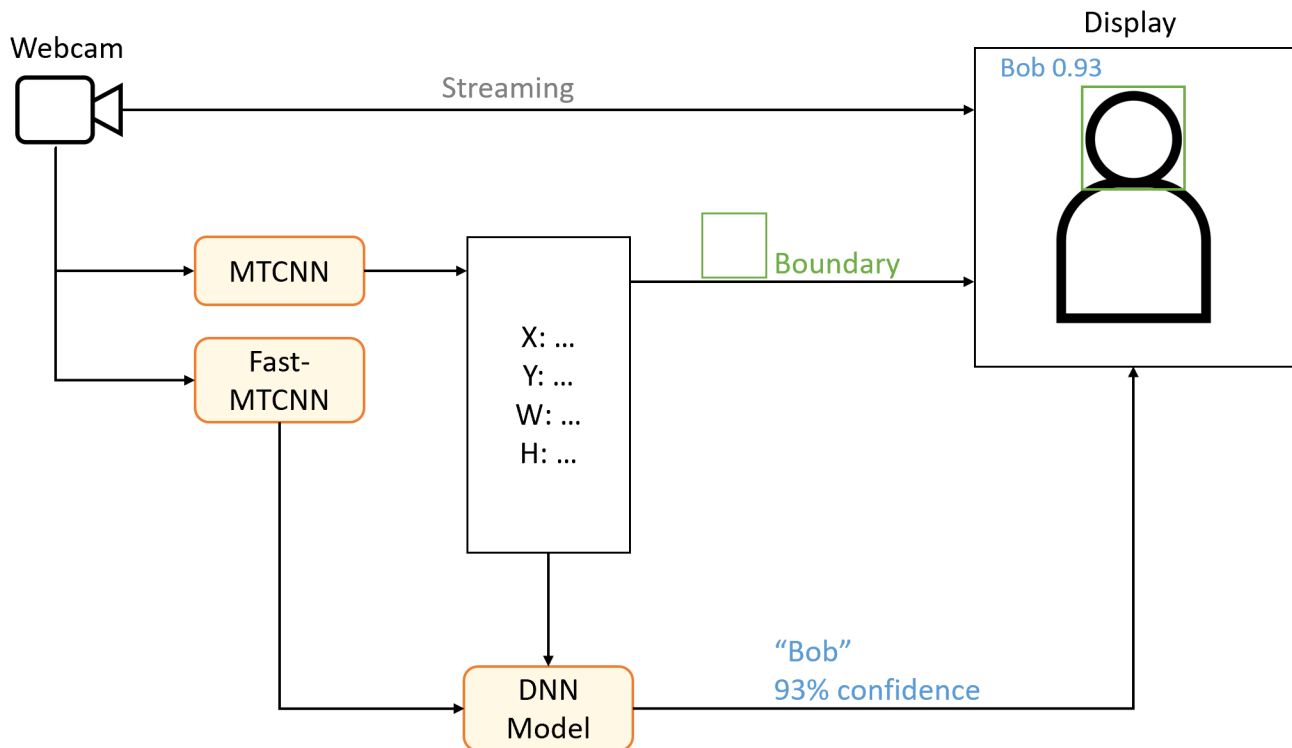
Human facial recognition is a classical task of computer vision. It has important applications in access control, attendance counting, facial security checks, etc. With the outbreak of the Covid-19, wearing masks has become mandatory for most public areas, which pose challenges for conventional facial recognition solutions. To solve this issue, our team plans to use a CNN-based deep learning model to achieve the masked face identity recognition task.

Our current progress is two-folds:

- We employed [MTCNN \(https://github.com/timesler/facenet-pytorch/blob/master/examples/infer.ipynb\)](https://github.com/timesler/facenet-pytorch/blob/master/examples/infer.ipynb) -- a facial detection framework, to crop the target facial area. Facial detection in video streams are supported, which enables testing in real-time.
- We use Alexnet as a pretrained network followed by a simple CNN to perform the (unmasked) human facial recognition task. The dataset we use for now is [the ORL dataset \(https://www.kaggle.com/tavarez/the-orl-database-for-training-and-testing\)](https://www.kaggle.com/tavarez/the-orl-database-for-training-and-testing).

Our plan for the next stage:

- In the next stage, we plan to crop the eye and forehead area with MTCNN and use this to train our neural network to do the masked facial recognition task. The most important idea is to exclude mask in the input feature so as not to let the neural network to learn the shape of the mask instead of the user identity.
- We also notice there are some traditional computer vision algorithms that could preprocess (e.g. Autolevel) the image or perform facial recognition tasks (e.g. local binary pattern). Therefore, we are also interested in whether incorporating these methods into our DNN could improve the model performance.



Part 1: Extract facial pictures from webcam streaming for ID recognition

Install necessary packages

```
In [1]: %%shell  
pip install facenet-pytorch  
pip install mmcv
```

Collecting facenet-pytorch

Downloading facenet_pytorch-2.5.2-py3-none-any.whl (1.9 MB)

|██| 1.9 MB 11.2 MB/s

Requirement already satisfied: pillow in /usr/local/lib/python3.7/dist-packages (from facenet-pytorch) (7.1.2)

Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from facenet-pytorch) (1.21.5)

Requirement already satisfied: torchvision in /usr/local/lib/python3.7/dist-packages (from facenet-pytorch) (0.11.1+cu111)

Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (from facenet-pytorch) (2.23.0)

Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests->facenet-pytorch) (3.0.4)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests->facenet-pytorch) (2021.10.8)

Requirement already satisfied: urllib3!=1.25.0,!>=1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from requests->facenet-pytorch) (1.24.3)

Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests->facenet-pytorch) (2.10)

Requirement already satisfied: torch==1.10.0 in /usr/local/lib/python3.7/dist-packages (from torchvision->facenet-pytorch) (1.10.0+cu111)

Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from torch==1.10.0->torchvision->facenet-pytorch) (3.10.0.2)

Installing collected packages: facenet-pytorch

Successfully installed facenet-pytorch-2.5.2

Collecting mmcv

Downloading mmcv-1.4.6.tar.gz (438 kB)

|██| 438 kB 6.4 MB/s

Collecting addict

Downloading addict-2.4.0-py3-none-any.whl (3.8 kB)

Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from mmcv) (1.21.5)

Requirement already satisfied: packaging in /usr/local/lib/python3.7/dist-packages (from mmcv) (21.3)

Requirement already satisfied: Pillow in /usr/local/lib/python3.7/dist-packages (from mmcv) (7.1.2)

Requirement already satisfied: pyyaml in /usr/local/lib/python3.7/dist-packages (from mmcv) (3.13)

Collecting yapf

Downloading yapf-0.32.0-py2.py3-none-any.whl (190 kB)

|██| 190 kB 22.3 MB/s

Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /usr/local/lib/python3.7/dist-packages (from packaging->mmcv) (3.0.7)

Building wheels for collected packages: mmcv

Building wheel for mmcv (setup.py) ... done

Created wheel for mmcv: filename=mmcv-1.4.6-py2.py3-none-any.whl size=656045 sha256=a0a1c1f7a98b3a97f923d46ea0959ded6a0fec77aabf3648e91f3afcb839bda0

Stored in directory: /root/.cache/pip/wheels/4b/dc/28/f5fdb35b7e1a5f50de1a95a49e5f661e4ffb10461d35974240

Successfully built mmcv

Installing collected packages: yapf, addict, mmcv

Successfully installed addict-2.4.0 mmcv-1.4.6 yapf-0.32.0

Out[1]:

Import Dependencies

```
In [2]: from IPython.display import Javascript, Image
        from IPython import display as dis
        from google.colab.output import eval_js
        from base64 import b64decode, b64encode
        import numpy as np
        import io
        import html
        import time
        from IPython.core.display import Video
        from facenet_pytorch import MTCNN
        import torch
        import mmcv, cv2
        import PIL
        from PIL import Image, ImageDraw
        import matplotlib.pyplot as plt
```

Several Helper methods to create video streaming

```

In [3]: # function to convert the JavaScript object into an OpenCV image
def js_to_image(js_reply):
    image_bytes = b64decode(js_reply.split(',')[1]) # decode base64 image
    jpg_as_np = np.frombuffer(image_bytes, dtype=np.uint8) # convert bytes to numpy array
    img = cv2.imdecode(jpg_as_np, flags=1) # decode numpy array into OpenCV BGR image

    return img

# function to convert OpenCV Rectangle bounding box image into base64 byte string to be overlaid on video stream
def bbox_to_bytes(ipt):
    # ipt: Numpy array (pixels) containing rectangle to overlay on video stream.
    # bytes: Base64 image byte string
    # convert array into PIL image
    bbox_PIL = PIL.Image.fromarray(ipt, 'RGBA')
    iobuf = io.BytesIO()
    # format bbox into png for return
    bbox_PIL.save(iobuf, format='png')
    # format return string
    bbox_bytes = 'data:image/png;base64,{}'.format((str(b64encode(iobuf.getvalue()))), 'utf-8'))

    return bbox_bytes

# JavaScript to properly create our live video stream using our webcam as input
def video_stream():
    js = Javascript('''
    var video;
    var div = null;
    var stream;
    var captureCanvas;
    var imgElement;
    var labelElement;
    var pendingResolve = null;
    var shutdown = false;

    function removeDom() {
        stream.getVideoTracks()[0].stop();
        video.remove();
        div.remove();
        video = null;
        div = null;
        stream = null;
        imgElement = null;
        captureCanvas = null;
        labelElement = null;
    }

    function onAnimationFrame() {
        if (!shutdown) {
            window.requestAnimationFrame(onAnimationFrame);
        }
        if (pendingResolve) {

```

```

        var result = "";
        if (!shutdown) {
            captureCanvas.getContext('2d').drawImage(video, 0, 0, 720, 720);
            result = captureCanvas.toDataURL('image/jpeg', 0.8)
        }
        var lp = pendingResolve;
        pendingResolve = null;
        lp(result);
    }
}

async function createDom() {
    if (div !== null) {
        return stream;
    }

    div = document.createElement('div');
    div.style.border = '2px solid black';
    div.style.padding = '3px';
    div.style.width = '100%';
    div.style.maxWidth = '600px';
    document.body.appendChild(div);

    const modelOut = document.createElement('div');
    modelOut.innerHTML = "<span>Status:</span>";
    labelElement = document.createElement('span');
    labelElement.innerText = 'No data';
    labelElement.style.fontWeight = 'bold';
    modelOut.appendChild(labelElement);
    div.appendChild(modelOut);

    video = document.createElement('video');
    video.style.display = 'block';
    video.width = div.clientWidth - 6;
    video.setAttribute('playsinline', '');
    video.onclick = () => { shutdown = true; };
    stream = await navigator.mediaDevices.getUserMedia(
        {video: { facingMode: "environment"}});
    div.appendChild(video);

    imgElement = document.createElement('img');
    imgElement.style.position = 'absolute';
    imgElement.style.zIndex = 1;
    imgElement.onclick = () => { shutdown = true; };
    div.appendChild(imgElement);

    const instruction = document.createElement('div');
    instruction.innerHTML = '<span style="color: red; font-weight: bold;">'
+ 'When finished, click here or on the video to stop this demo</span>';
    div.appendChild(instruction);
    instruction.onclick = () => { shutdown = true; };

    video.srcObject = stream;
    await video.play();

    captureCanvas = document.createElement('canvas');
    captureCanvas.width = 720; //video width; 1280

```

```

captureCanvas.height = 720; //video height; 720
window.requestAnimationFrame(onAnimationFrame);

return stream;
}
async function stream_frame(label, imgData) {
  if (shutdown) {
    removeDom();
    shutdown = false;
    return '';
  }

  var preCreate = Date.now();
  stream = await createDom();

  var preShow = Date.now();
  if (label !== "") {
    labelElement.innerHTML = label;
  }

  if (imgData !== "") {
    var videoRect = video.getClientRects()[0];
    imgElement.style.top = videoRect.top + "px";
    imgElement.style.left = videoRect.left + "px";
    imgElement.style.width = videoRect.width + "px";
    imgElement.style.height = videoRect.height + "px";
    imgElement.src = imgData;
  }

  var preCapture = Date.now();
  var result = await new Promise(function(resolve, reject) {
    pendingResolve = resolve;
  });
  shutdown = false;

  return {'create': preShow - preCreate,
        'show': preCapture - preShow,
        'capture': Date.now() - preCapture,
        'img': result};
}
'''

display(js)

def video_frame(label, bbox):
  data = eval_js('stream_frame("{}","{}").format(label, bbox))
  return data

```

Constructing a faster version of face detection model for potential higher throughput demand


```
In [4]: class FastMTCNN(object):
        """Fast MTCNN implementation."""

        def __init__(self, stride, resize=1, *args, **kwargs):
            """Constructor for FastMTCNN class.

            Arguments:
                stride (int): The detection stride. Faces will be detected every `
stride` frames
                        and remembered for `stride-1` frames.

            Keyword arguments:
                resize (float): Fractional frame scaling. [default: {1}]
                *args: Arguments to pass to the MTCNN constructor. See help(MTCN
N).
                **kwargs: Keyword arguments to pass to the MTCNN constructor. See
help(MTCNN).
            """
            self.stride = stride
            self.resize = resize
            self.mtcnn = MTCNN(*args, **kwargs)

        def __call__(self, frames):
            """Detect faces in frames using strided MTCNN."""
            if self.resize != 1:
                frames = [
                    cv2.resize(f, (int(f.shape[1] * self.resize), int(f.shape[0] *
self.resize)))
                    for f in frames
                ]

            boxes, probs = self.mtcnn.detect(frames[::self.stride])

            faces = []
            for i, frame in enumerate(frames):
                box_ind = int(i / self.stride)
                if boxes[box_ind] is None:
                    continue
                for box in boxes[box_ind]:
                    box = [int(b) for b in box]
                    faces.append(frame[box[1]:box[3], box[0]:box[2]])

            return faces
```

Face detection model implementations

```
In [5]: device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print('Running on device: {}'.format(device))
# currently using, for face detection
mtcnn = MTCNN(keep_all=True, min_face_size=224, device=device)
# A faster model
mtcnn_fast = FastMTCNN(keep_all=True, min_face_size=224, device=device, stride=
4)
```

Running on device: cuda:0

```

In [14]: from torch.functional import Tensor
video_stream()
label_html = 'Capturing...'
bbox = ''
while True:
    js_reply = video_frame(label_html, bbox)
    if not js_reply:
        break
    # convert JS response to OpenCV Image
    img = js_to_image(js_reply["img"])

    rgb_img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    #frame_fast = torch.from_numpy(rgb_img)
    #frame_fast = torch.stack([frame_fast], dim=0) #Prepared input frame for m
tcnn_fast model
    frame = Image.fromarray(rgb_img) #Input frame for mtcnn face detection mod
el

    faces, _ = mtcnn.detect(frame) #return coordinates of faces
    #faces_fast = mtcnn_fast(frame_fast) #higher throughput, return captured f
aces
    #plt.imshow(faces_fast[0].detach().numpy(), cmap = 'gray')
    #plt.show()

    bbox_array = np.zeros([720,720,4], dtype=np.uint8) #Transparent overlay fo
r drawing bounding box

    if faces is not None:
        for (x,y,w,h) in faces:
            x_w_diff = int(w-x)
            y_h_diff = int(h-y)
            if x_w_diff > 223. and y_h_diff >223.:
                x_w_mid = int(x+(w-x)/2)
                y_h_mid = int(y+(h-y)/2)
                selected_x = x_w_mid - 112
                selected_w = x_w_mid + 112
                selected_y = y_h_mid - 112 - 50
                selected_h = y_h_mid + 112 - 50

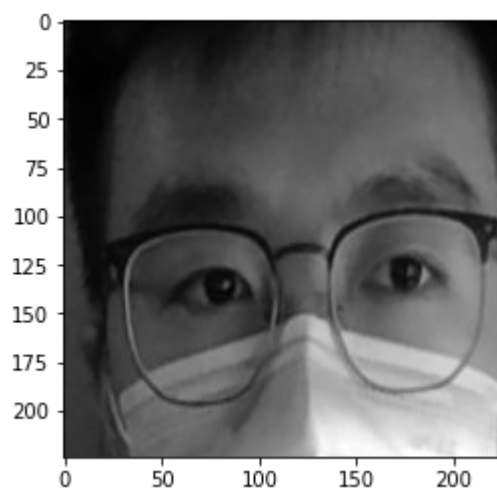

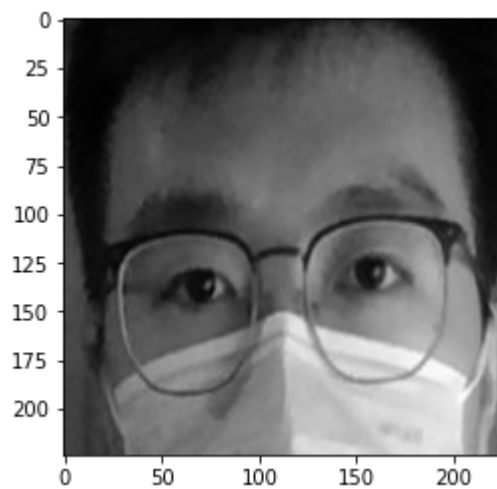
                #selected_img = img[int(y):int(y) + x_w_diff,int(x):int(x)+ x_w_dif
f,: ]
                selected_img = img[selected_y:selected_h, selected_x:selected_w, :]
                selected_frame = cv2.cvtColor(selected_img, cv2.COLOR_BGR2GRAY)
                plt.imshow(selected_frame, cmap = 'gray')
                plt.show()
                print(selected_frame.shape)


                input_for_face_recognition = torch.from_numpy(selected_frame)
                input_for_face_recognition = torch.stack([input_for_face_recognition
], dim=0) #Prepared input in tensor format

                # Facial recognition process will be placed here

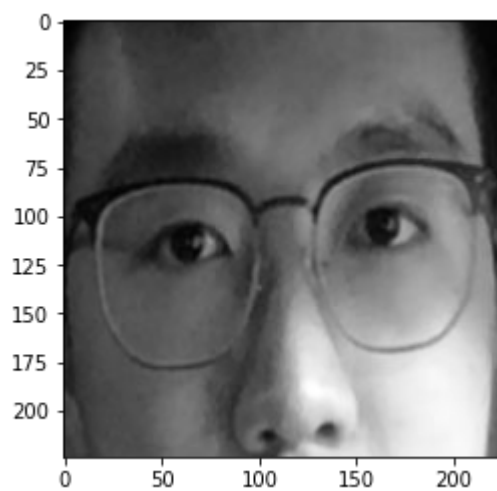
```

```
#
box_height = (w,int(h-(20)))
bbox_array = cv2.rectangle(bbox_array,(x,y), box_height, (0,255,0),
2) #Bounding box size and colour
bbox_array[:, :, 3] = (bbox_array.max(axis = 2) > 0 ).astype(int) * 255
bbox_bytes = bbox_to_bytes(bbox_array) # converting overlay of bbox into
bytes
bbox =bbox_bytes # update bbox for the next frame
```





(224, 224)



(224, 224)

```
In [ ]: %%shell
jupyter nbconvert --to html Project_Face_detection.ipynb
```

Part 2: Model Training

```
In [1]: from google.colab import drive  
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call `drive.mount("/content/drive", force_remount=True)`.

```
In [2]: from os import listdir  
from os.path import isfile, join  
import os  
import re  
import shutil  
import numpy as np  
import torch  
import torch.nn as nn  
import torch.nn.functional as F  
import torchvision  
from torch.utils.data.sampler import SubsetRandomSampler  
from torchvision import datasets, transforms  
import torch.optim as optim  
import matplotlib.pyplot as plt  
from tqdm import tqdm  
from PIL import Image  
import matplotlib.pyplot as plt  
  
use_cuda = True
```

```

In [3]: def split_data(dataset_path, batch_size = 64):
    torch.manual_seed(1)
    np.random.seed(1)

    transform = transforms.Compose([transforms.Resize((224,224)),
                                    # transforms.ColorJitter(), # change image color
                                    transforms.RandomHorizontalFlip(), # flip images
                                    # transforms.RandomAffine(30), # Random affine transformation of the image keeping center invariant.
                                    # transforms.CenterCrop(224),
                                    transforms.ToTensor()])

    training_dataset = datasets.ImageFolder(root = dataset_path, transform=transform)

    indices = list(range(len(training_dataset)))

    train_indices = []

    val_indices = []

    # split data with ratio 8:2
    for i in indices:
        if i % 10 > 7:
            val_indices.append(i)
        else:
            train_indices.append(i)

    np.random.shuffle(train_indices)

    np.random.shuffle(val_indices)

    train_sampler = SubsetRandomSampler(train_indices)
    val_sampler = SubsetRandomSampler(val_indices)

    train_loader, val_loader = torch.utils.data.DataLoader(training_dataset, batch_size=batch_size, sampler=train_sampler), \
                                torch.utils.data.DataLoader(training_dataset, batch_size=batch_size, sampler=val_sampler)

    return train_loader, val_loader

```

```

In [4]: # set data path
ori_path = '/content/drive/MyDrive/ut/MIE1517_project_dataset/orl/images/'
group_folder_path = '/content/drive/MyDrive/ut/MIE1517_project_dataset/orl/grouped_images/'
files = [f for f in listdir(ori_path) if isfile(join(ori_path, f))]

```



```
In [5]: # extract images and labels and store into target folder
for single_file in files:
    group = single_file.split("_")[1].split('.')[0]
    folder_path = group_folder_path + group

    if not os.path.exists(folder_path):
        print("Directory '{}' is created!".format(group))
        os.makedirs(folder_path)

    shutil.copyfile(ori_path + single_file, folder_path + '/' + single_file)
```

```
In [6]: group_folder_path = '/content/drive/MyDrive/ut/MIE1517_project_dataset/orl/grouped_images/'

train_loader, val_loader = split_data(group_folder_path, batch_size = 1)
```

```
In [7]: print('Training images we have:', len(train_loader))
        print('Validation images we have:', len(val_loader))
```

Training images we have: 328
Validation images we have: 82

```
In [8]: # define model to adapt Alex Model
class AlexNetModel(nn.Module):

    def __init__(self):
        super(AlexNetModel, self).__init__()
        self.name = 'AlexNetModel'
        self.conv1 = nn.Conv2d(256, 30, 3)
        self.pool = nn.MaxPool2d(2, 2)
        self.fc1 = nn.Linear(120, 64)
        self.fc2 = nn.Linear(64, 41)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = x.view(-1, 120)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x
```

```

In [11]: # function to save checkpoint
def save_checkpoint(model, batch_size, lr, epoch):
    model_path = "{}_{}_{}_{}".format(model.name, batch_size, lr, epoch)
    torch.save(model, model_path)
    print('Checkpoint of {} has been stored successfully!',format(model_path))

def get_accuracy(model, data_loader, grey_images_flag = False):
    correct = 0
    total = 0

    alexnet = torchvision.models.alexnet(pretrained=True)

    for imgs, labels in data_loader:
        #####
        #To Enable GPU Usage

        if grey_images_flag:

            grey_images = torchvision.transforms.Grayscale()(imgs)

            imgs = torch.tensor(np.tile(grey_images, [1,3,1,1]))

            imgs = alexnet.features(imgs)

            if use_cuda and torch.cuda.is_available():
                imgs = imgs.cuda()
                labels = labels.cuda()
            #####

            output = model(imgs)

            #select index with maximum prediction score
            pred = output.max(1, keepdim=True)[1]
            correct += pred.eq(labels.view_as(pred)).sum().item()
            total += imgs.shape[0]
    return correct / total

def train(model, dataset_path, opt, batch_size = 64, learning_rate=0.01, epoch
s=30, grey_images_flag = False):

    #####
    # input: grey_images_flag: boolean if gray image conversion is needed
    #####

    train_loader, val_loader = split_data(dataset_path, batch_size = batch_size)

    criterion = nn.CrossEntropyLoss()

    optimizer = opt(model.parameters(), lr=learning_rate)

    iters, losses, train_acc, val_acc = [], [], [], []

    alexnet = torchvision.models.alexnet(pretrained=True)

    n = 0 # the number of iterations
    for epoch in range(epochs):

```

```

for imgs, labels in tqdm(iter(train_loader)):

    if grey_images_flag:

        grey_images = torchvision.transforms.Grayscale()(imgs)

        imgs = torch.tensor(np.tile(grey_images, [1,3,1,1]))

    imgs = alexnet.features(imgs)

    #####
    #To Enable GPU Usage
    if use_cuda and torch.cuda.is_available():
        imgs = imgs.cuda()
        labels = labels.cuda()
    #####

    output = model(imgs)
    loss = criterion(output, labels)
    loss.backward()
    optimizer.step()
    optimizer.zero_grad()

    # save the current training information
    iters.append(n)
    losses.append(float(loss)/batch_size)
    train_acc.append(get_accuracy(model, train_loader, grey_images_flag = grey_images_flag))
    val_acc.append(get_accuracy(model, val_loader, grey_images_flag = grey_images_flag))
    n += 1

    if ((epoch+1) % 25 == 0):
        save_checkpoint(model, batch_size, learning_rate, epoch+1)

    # plotting
    plt.title("Training Curve")
    plt.plot(iters, losses, label="Train")
    plt.xlabel("Iterations")
    plt.ylabel("Loss")
    plt.show()

    plt.title("Training Curve")
    plt.plot(iters, train_acc, label="Train")
    plt.plot(iters, val_acc, label="Validation")
    plt.xlabel("Iterations")
    plt.ylabel("Training Accuracy")
    plt.legend(loc='best')
    plt.show()

    print("Final Training Accuracy: {}".format(train_acc[-1]))
    print("Final Validation Accuracy: {}".format(val_acc[-1]))

```

```
In [12]: model = AlexNetModel()

if use_cuda and torch.cuda.is_available():
    model.cuda()
    print('CUDA is available!  Training on GPU ...')
else:
    print('CUDA is not available.  Training on CPU ...')

train(model, group_folder_path, opt = optim.Adam, batch_size = 32, learning_rate=0.001, epochs=100, grey_images_flag = False)
```

CUDA is not available. Training on CPU ...

100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.16s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:24<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:24<00:00,	2.18s/it]
100%		11/11	[00:24<00:00,	2.19s/it]
100%		11/11	[00:24<00:00,	2.20s/it]
100%		11/11	[00:24<00:00,	2.19s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:24<00:00,	2.19s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:24<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.18s/it]

Checkpoint of {} has been stored successfully! AlexNetModel_32_0.001_25

100%		11/11	[00:24<00:00,	2.19s/it]
100%		11/11	[00:24<00:00,	2.18s/it]
100%		11/11	[00:24<00:00,	2.20s/it]
100%		11/11	[00:24<00:00,	2.19s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:24<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:24<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:24<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:24<00:00,	2.19s/it]
100%		11/11	[00:24<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.18s/it]

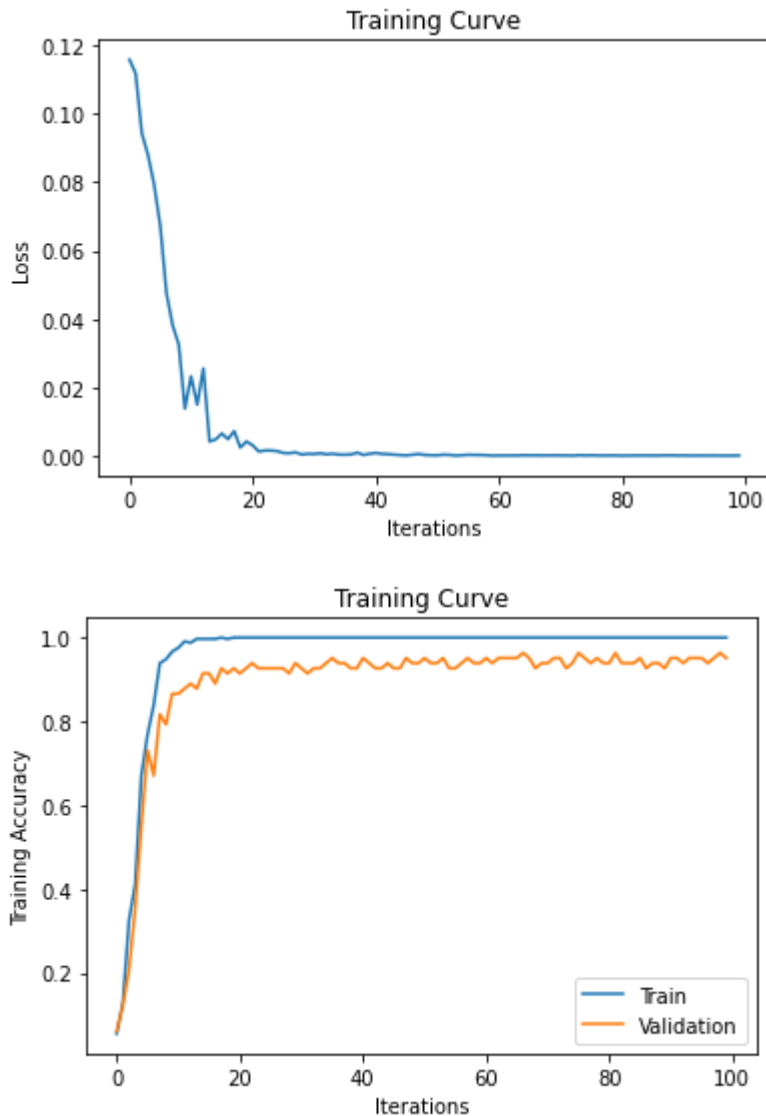
Checkpoint of {} has been stored successfully! AlexNetModel_32_0.001_50

100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.16s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:24<00:00,	2.19s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.16s/it]
100%		11/11	[00:24<00:00,	2.19s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.18s/it]

Checkpoint of {} has been stored successfully! AlexNetModel_32_0.001_75

100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:24<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.16s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.16s/it]
100%		11/11	[00:23<00:00,	2.16s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.17s/it]
100%		11/11	[00:24<00:00,	2.18s/it]
100%		11/11	[00:23<00:00,	2.18s/it]
100%		11/11	[00:24<00:00,	2.18s/it]
100%		11/11	[00:24<00:00,	2.19s/it]
100%		11/11	[00:24<00:00,	2.19s/it]

Checkpoint of {} has been stored successfully! AlexNetModel_32_0.001_100



Final Training Accuracy: 1.0

Final Validation Accuracy: 0.9512195121951219

Future plan

As training and testing dataset is clean and standard, adding noise is not necessary for current phase but could be benefit to real world data. Also, image processing method such as auto-level to scale the distribution of pixels for each channel could also improve our model performance for real world data.

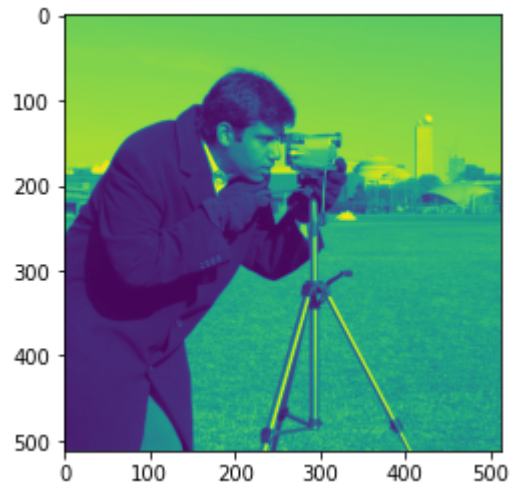
Example:

```
In [20]: from skimage import data
from skimage.morphology import disk
from skimage.filters.rank import autolevel
from PIL import Image
```

```
In [16]: img = data.camera()  
auto = autolevel(img, disk(5))
```

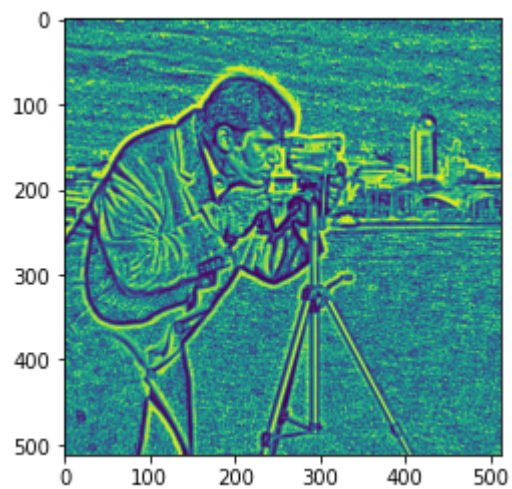
```
In [17]: plt.imshow(img)
```

Out[17]: <matplotlib.image.AxesImage at 0x7f5059ed0950>



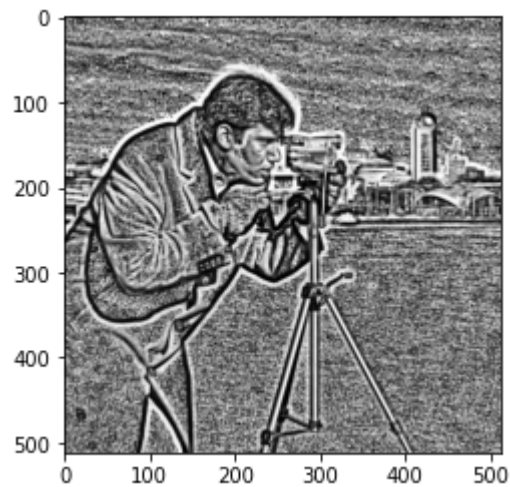
```
In [18]: plt.imshow(auto)
```

Out[18]: <matplotlib.image.AxesImage at 0x7f50585b2a90>





```
In [19]: plt.imshow(auto, cmap = "gray")
```

```
Out[19]: <matplotlib.image.AxesImage at 0x7f505851ff10>
```



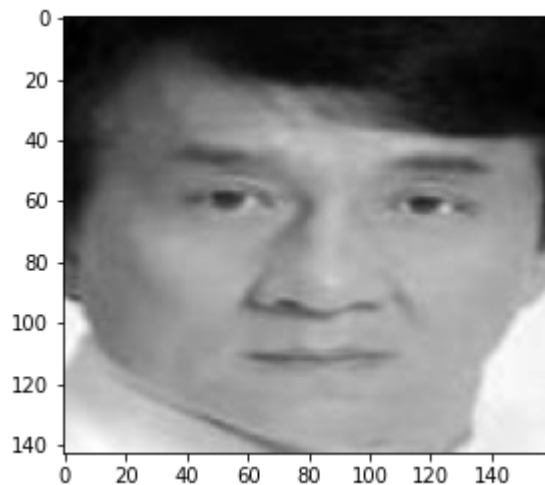
```
In [40]: image = Image.open('img_1.jpg')  
plt.imshow(image)
```

```
Out[40]: <matplotlib.image.AxesImage at 0x7f5058021990>
```



```
In [46]: image = torchvision.transforms.Grayscale()(image)
plt.imshow(image, cmap='gray')
```


Out[46]: <matplotlib.image.AxesImage at 0x7f5057f14ad0>



```
In [59]: image = np.asarray(image)
image_new = image.copy()
```

```
In [65]: plt.imshow(autolevel(image_new, disk(5)), cmap='gray')
```

Out[65]: <matplotlib.image.AxesImage at 0x7f5057c1c1d0>



```
In [67]: def auto_level_img(imgs):
ret_imgs = []
imgs = torchvision.transforms.Grayscale()(imgs)
for i in imgs:
ret_imgs.append(autolevel(image_new, disk(5)))
```

```

In [68]: def train_autolevel(model, dataset_path, opt, batch_size = 64, learning_rate=
0.01, epochs=30, grey_images_flag = False):

    #####
    # input: grey_images_flag: boolean if gray image conversion is needed
    #####

    train_loader, val_loader = split_data(dataset_path, batch_size = batch_size)

    criterion = nn.CrossEntropyLoss()

    optimizer = opt(model.parameters(), lr=learning_rate)

    iters, losses, train_acc, val_acc = [], [], [], []

    alexnet = torchvision.models.alexnet(pretrained=True)

    n = 0 # the number of iterations
    for epoch in range(epochs):
        for imgs, labels in tqdm(iter(train_loader)):

            if grey_images_flag:

                grey_images = torchvision.transforms.Grayscale()(imgs)

                grey_images = auto_level_img(grey_images)

                imgs = torch.tensor(np.tile(grey_images, [1,3,1,1]))

            imgs = alexnet.features(imgs)

            #####
            #To Enable GPU Usage
            if use_cuda and torch.cuda.is_available():
                imgs = imgs.cuda()
                labels = labels.cuda()
            #####

            output = model(imgs)
            loss = criterion(output, labels)
            loss.backward()
            optimizer.step()
            optimizer.zero_grad()

            # save the current training information
            iters.append(n)
            losses.append(float(loss)/batch_size)
            train_acc.append(get_accuracy(model, train_loader, grey_images_flag = grey
_images_flag))
            val_acc.append(get_accuracy(model, val_loader, grey_images_flag = grey_ima
ges_flag))
            n += 1

            if ((epoch+1) % 25 == 0):
                save_checkpoint(model, batch_size, learning_rate, epoch+1)

```

```
# plotting
plt.title("Training Curve")
plt.plot(iters, losses, label="Train")
plt.xlabel("Iterations")
plt.ylabel("Loss")
plt.show()

plt.title("Training Curve")
plt.plot(iters, train_acc, label="Train")
plt.plot(iters, val_acc, label="Validation")
plt.xlabel("Iterations")
plt.ylabel("Training Accuracy")
plt.legend(loc='best')
plt.show()

print("Final Training Accuracy: {}".format(train_acc[-1]))
print("Final Validation Accuracy: {}".format(val_acc[-1]))
```

```
In [69]: model_auto_level = AlexNetModel()

if use_cuda and torch.cuda.is_available():
    model_auto_level.cuda()
    print('CUDA is available! Training on GPU ...')
else:
    print('CUDA is not available. Training on CPU ...')

train(model_auto_level, group_folder_path, opt = optim.Adam, batch_size = 32,
learning_rate=0.001, epochs=100, grey_images_flag = True)
```

CUDA is not available. Training on CPU ...

100%	██████████	11/11	[00:25<00:00,	2.32s/it]
100%	██████████	11/11	[00:24<00:00,	2.20s/it]
100%	██████████	11/11	[00:24<00:00,	2.19s/it]
100%	██████████	11/11	[00:24<00:00,	2.19s/it]
100%	██████████	11/11	[00:24<00:00,	2.20s/it]
100%	██████████	11/11	[00:24<00:00,	2.21s/it]
100%	██████████	11/11	[00:24<00:00,	2.20s/it]
100%	██████████	11/11	[00:24<00:00,	2.20s/it]
100%	██████████	11/11	[00:24<00:00,	2.22s/it]
100%	██████████	11/11	[00:24<00:00,	2.21s/it]
100%	██████████	11/11	[00:24<00:00,	2.22s/it]
100%	██████████	11/11	[00:24<00:00,	2.22s/it]
100%	██████████	11/11	[00:24<00:00,	2.22s/it]
100%	██████████	11/11	[00:24<00:00,	2.23s/it]
100%	██████████	11/11	[00:24<00:00,	2.22s/it]
100%	██████████	11/11	[00:24<00:00,	2.21s/it]
100%	██████████	11/11	[00:24<00:00,	2.21s/it]
100%	██████████	11/11	[00:24<00:00,	2.21s/it]
100%	██████████	11/11	[00:24<00:00,	2.20s/it]
100%	██████████	11/11	[00:24<00:00,	2.21s/it]
100%	██████████	11/11	[00:24<00:00,	2.20s/it]
100%	██████████	11/11	[00:24<00:00,	2.21s/it]
100%	██████████	11/11	[00:24<00:00,	2.21s/it]
100%	██████████	11/11	[00:24<00:00,	2.20s/it]
100%	██████████	11/11	[00:24<00:00,	2.20s/it]

Checkpoint of {} has been stored successfully! AlexNetModel_32_0.001_25

100%	██████████	11/11	[00:24<00:00,	2.20s/it]
100%	██████████	11/11	[00:24<00:00,	2.19s/it]
100%	██████████	11/11	[00:24<00:00,	2.19s/it]
100%	██████████	11/11	[00:24<00:00,	2.19s/it]
100%	██████████	11/11	[00:24<00:00,	2.19s/it]
100%	██████████	11/11	[00:24<00:00,	2.21s/it]
100%	██████████	11/11	[00:24<00:00,	2.20s/it]
100%	██████████	11/11	[00:24<00:00,	2.19s/it]
100%	██████████	11/11	[00:24<00:00,	2.25s/it]
100%	██████████	11/11	[00:23<00:00,	2.18s/it]
100%	██████████	11/11	[00:24<00:00,	2.19s/it]
100%	██████████	11/11	[00:23<00:00,	2.18s/it]
100%	██████████	11/11	[00:24<00:00,	2.19s/it]
100%	██████████	11/11	[00:24<00:00,	2.20s/it]
100%	██████████	11/11	[00:24<00:00,	2.19s/it]
100%	██████████	11/11	[00:24<00:00,	2.20s/it]
100%	██████████	11/11	[00:24<00:00,	2.21s/it]
100%	██████████	11/11	[00:24<00:00,	2.20s/it]
100%	██████████	11/11	[00:24<00:00,	2.21s/it]
100%	██████████	11/11	[00:24<00:00,	2.20s/it]
100%	██████████	11/11	[00:24<00:00,	2.20s/it]
100%	██████████	11/11	[00:24<00:00,	2.20s/it]
100%	██████████	11/11	[00:24<00:00,	2.19s/it]
100%	██████████	11/11	[00:24<00:00,	2.21s/it]
100%	██████████	11/11	[00:24<00:00,	2.19s/it]



Checkpoint of {} has been stored successfully! AlexNetModel_32_0.001_50

100%		11/11	[00:24<00:00,	2.20s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.21s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.21s/it]
100%		11/11	[00:24<00:00,	2.21s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.23s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.21s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.21s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.23s/it]
100%		11/11	[00:24<00:00,	2.21s/it]
100%		11/11	[00:24<00:00,	2.22s/it]

Checkpoint of {} has been stored successfully! AlexNetModel_32_0.001_75

100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.21s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.21s/it]
100%		11/11	[00:24<00:00,	2.21s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.23s/it]
100%		11/11	[00:24<00:00,	2.23s/it]
100%		11/11	[00:24<00:00,	2.21s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.23s/it]
100%		11/11	[00:24<00:00,	2.23s/it]
100%		11/11	[00:24<00:00,	2.24s/it]
100%		11/11	[00:24<00:00,	2.23s/it]
100%		11/11	[00:24<00:00,	2.24s/it]
100%		11/11	[00:24<00:00,	2.22s/it]
100%		11/11	[00:24<00:00,	2.24s/it]

Checkpoint of {} has been stored successfully! AlexNetModel_32_0.001_100



Final Training Accuracy: 1.0

Final Validation Accuracy: 0.9878048780487805

In []: